

Common Quiz Mistakes





UVA CS

Machine Learning Fall 2020

Quiz 3

Full Gradient Descent can be viewed as a special case of Mini-batch SGD


27 / 39 correct responses



Quiz 3

Which data splits should you train on?


39 / 39 correct responses



Quiz 3

If you are given N data points and use half for training and the other for testing, the difference between training error and test error decreases as N increases.

31 / 39 correct responses




Quiz 5

In k-nearest neighbor classifiers, as the number of training samples N increases, the time it takes to test a new sample:





30 / 40 correct responses



Quiz 6

At the decision boundary of logistic regression, $P(y = 1 | x) = P(y = 0 | x) =$



24 / 36 correct responses



Quiz 8

For the same input size, a convolution layer has more parameters to train than a fully connected (i.e. MLP) layer.


30 / 37 correct responses



Quiz 10

Which of the following are used to represent words so that semantics are preserved. For example, "king - man + woman = queen"


28 / 35 correct responses



Quiz 10

Which type of model is best for unordered input data (e.g. a set)

28 / 35 correct responses



Quiz 12

Compared to a single decision tree, random forests typically:

31 / 36 correct responses

